

PCSE 595

Special Topics in Machine Learning

Dr. Sam Henry

samuel.henry@cnu.edu

Luter 325

Office Hours

- Please stop by!
- This is an [Open Question Answering Time](#)

Monday, Wednesday, Friday
11:00-12:00, 1:00-1:30
Or by appointment

Office hours are in-person (just come by my office, LUTR 325)



*I expect to see
everyone at office
hours at some point*

Pizza My Mind

- You can get extra credit
 - Up to two extra points on your final grade for one PCSE course
- Attend them! They're fun, informative, and employers present
 - Don't wait until you need a job or internship, go now!
- Thursdays at 12:20

... and you get free pizza!!



Students in PCSE classes can get extra credit if they attend at least 10 events. 10-11 events: 1 extra point; 12-13 events: 2 extra points.

Data Balance

Breast Cancer Dataset

- 458 benign and 241 Malignant





Figure 1: Initial Approximate Boundaries of Cell Nuclei

Fairly Balanced Dataset
(~2 to 1 class balance)

Iris Dataset

- 50 of each flower type



Perfectly Balanced Dataset
(1 to 1 to 1 balance)

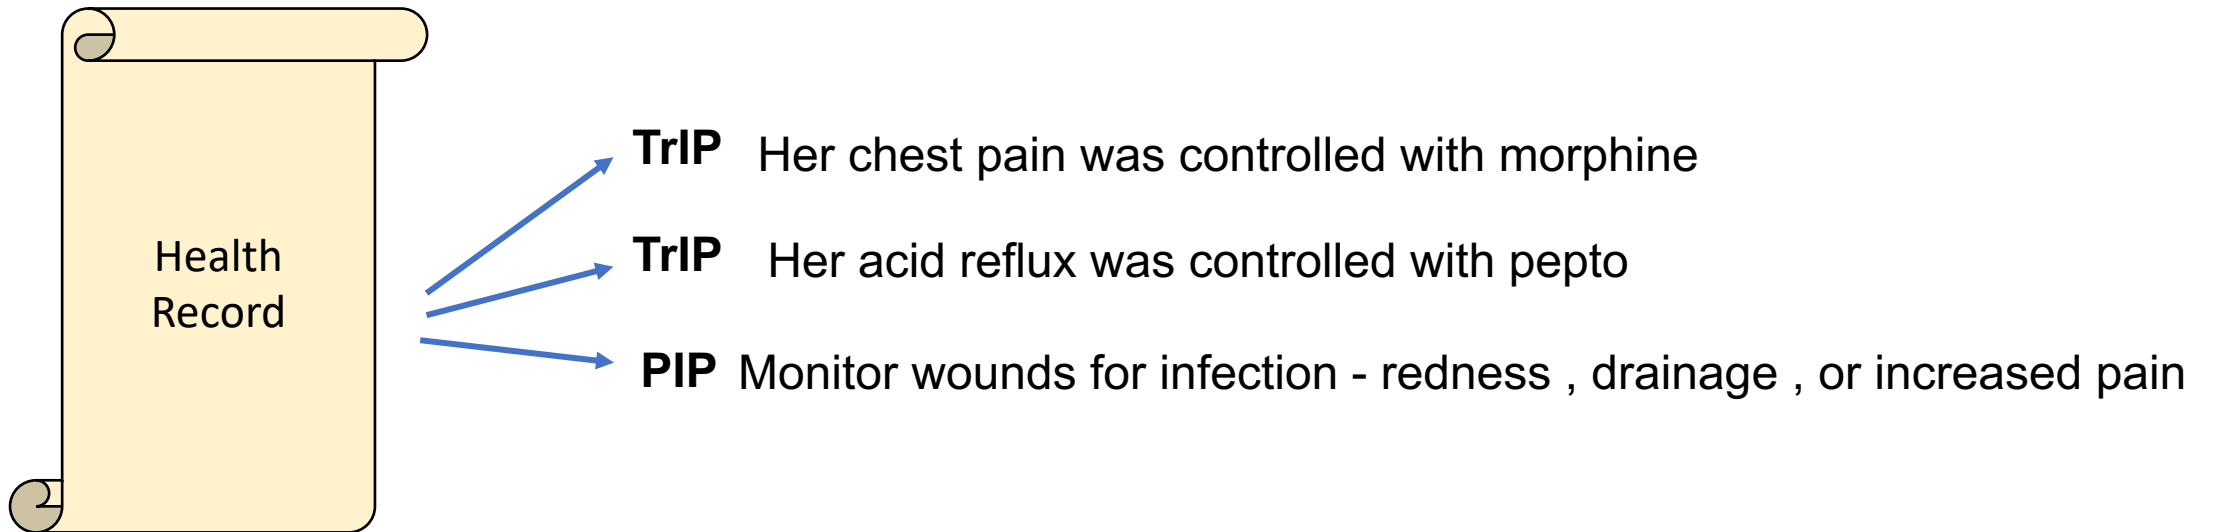
Consider the Following Problem

- Goal – Relationship Extraction
 - Given an electronic health record, extract all of the following relationships

Relationship	Definition	Example
TrIP	Treatment improves or cures medical problem	Her chest pain was controlled with morphine
TrWP	Treatment worsens or does not improve medical problem	... who presented with acute coronary syndrome refractory to medical treatment and TNK ...
TrCP	Treatment caused medical problem	Narcotics can cause constipation
TrAP	Treatment administered for medical problem, but outcome is not mentioned in the sentence	HTN elev. chol. right facial droop with metal plate secondary to GSW to face right nephrectomy
TrNAP	Treatment is not administered or discontinued because of a medical problem	Discharge Instructions : may shower , no bathing or swimming for 1 month no creams , lotions or powders to any incisions no driving for 1 month no lifting > 10 # for 10 weeks
TeRP	Test reveals a medical problem	Cath at Kindred/North Shore today showed 80% LM lesion with normal LAD , CX , RCA .
TeCP	Test given to investigate a medical problem	He also has a diagnosis of interstitial cystitis and more significantly has a history of coronary artery disease with a recent catheterization , supraventricular tachycardia , and a progressive mitochondrial myopathy .
PIP	Medical problem indicates or reveals aspects of another medical problem	Monitor wounds for infection - redness , drainage , or increased pain

Relationship Extraction

- Goal – Relationship Extraction
 - Given an electronic health record, extract relationships



Relationship Extraction as Sentence Classification

- Cast the problem as a sentence classification task
- For each sentence in the document, label it with each of the relations it contains

None Patient Admitted on 03/15/12

None Primary Complaints are chest pain and acid reflux

TrIP Her chest pain was controlled with morphine

TrIP Her acid reflux was controlled with pepto

TrCP During treatment, she was wounded

PIP Monitor wounds for infection - redness , drainage , or increased pain

Sentence Classification

- This is a Complicated Problem:

16,316 sentences

This is not binary classification – a sentence can contain 1 or more relationship types

Relationship	TrIP	TrWP	TrCP	TrAP	TrNAP	TeRP	TeCP	PIP	Total Relations	No Relation
Total Count	203	133	526	2617	174	3053	504	2203	9413	10,308
Unique Sentence Count	140	86	370	1643	119	1858	345	1447	6008	10,308

There is a huge class imbalance

$$\frac{6008}{86} \sim 70 \text{ to } 1 \text{ class ratio}$$

There are a lot of negative samples
If we consider “No Relation” a class,
then $\frac{10308}{86} \sim 120 \text{ to } 1 \text{ class ratio}$

Complex Problems

- Problems with multiple classes require:
 - Different classification architectures
 - Different evaluation metrics
 - Metrics for each class, metrics to combine results for all classes
 - More complex error analysis
- Imbalanced Datasets require:
 - Different evaluation measures
 - May have trouble training due to class imbalance
 - Results may need to be interpreted more carefully

Evaluation vs. Error Analysis

Evaluation

- Methods to quantify the performance of an algorithm

Error Analysis

- Methods to determine what needs to be improved to increase performance

Evaluation Metrics for Classifiers

- Primary Metrics:
 - Accuracy
 - Precision
 - Recall
 - F-measure
 - Receiver Operator Characteristic (ROC) Curves
 - Precision and Recall Curves
 - Correlation Coefficients

...and more

Why not just use accuracy?

Accuracy = Total Correct / Size of the dataset

90 samples of class 1
10 samples of class 2

Classify all as class 1

90/100 = 90% accuracy

- Probably the most popular measure
- Biased in favor of the majority class
 - assumes a 1 to 1 class distribution

Use with Caution!

Confusion Matrix and Types of Errors

confusion matrix

		Predicted Class		
		Class 1 (pos)	Class 2 (neg)	
True Class	Class 1 (pos)	TP	FN	P=TP+FN
	Class 2 (neg)	FP	TN	N=FP+TN

Type II Error

Type I Error

TP = True Positive = predicted positive and is positive

FP = False Positive = predicted positive but is negative

TN = True Negative = predicted negative and is negative

FN = False Negative = predicted negative but is positive

P = Number Positive = TP + FN

N = Number Negative = FP+TN

Accuracy

		Predicted Class		
		Class 1 (pos)	Class 2 (neg)	
True Class	Class 1 (pos)	TP	FN	P=TP+FN
	Class 2 (neg)	FP	TN	N=FP+TN

$$Accuracy = \frac{TP + TN}{P + N}$$

Measures the percent of predictions that were correct

Places an equal importance on positive and negative classes

Appropriate for balanced datasets

Range: (0,1); 1 is best

Classes aren't always the same

Accuracy assumes classes are balanced and are equally important

- In many cases, one class is more important than another
 - Fraud detection, cancer diagnosis, intrusion detection, relationship extraction
 - In these cases, we may tolerate greater overall error, in return for better predictions of the more important class
- Often the classes are imbalanced
 - In these cases, the performance on the **minority class** is hidden by the size of the **majority class**

Precision

		Predicted Class		
		Class 1 (pos)	Class 2 (neg)	
True Class	Class 1 (pos)	TP	FN	P=TP+FN
	Class 2 (neg)	FP	TN	N=FP+TN

$$Precision = \frac{TP}{TP + FP}$$

How good is my performance on the samples I am predicting to be true?

What percent of samples that I predict are true are actually true?

Range: (0,1); 1 is best

Recall

		Predicted Class		
		Class 1 (pos)	Class 2 (neg)	
True Class	Class 1 (pos)	TP	FN	P=TP+FN
	Class 2 (neg)	FP	TN	N=FP+TN


$$Recall = \frac{TP}{TP + FN}$$

What percentage of true samples am I predicting to be true?

Am I missing a lot of true samples?

Range: (0,1); 1 is best

Evaluation Example



There is a possible relationships
between each word and every
other word in the sentence

You predict:

aspirin, for
aspirin, and
aspirin, headache

You Miss:

Aspirin, neckpain
headache, aspirin
neckpain, aspirin

Calculate the confusion matrix, Accuracy, Precision, and Recall

What is the class ratio?

Precision and Recall Tradeoff

There is a trade-off between precision and recall

There is typically a threshold you can adjust to adjust the trade-off

High Recall = Fishing with a big net

High Precision = Spear fishing

- As you increase the size of your net, you catch more fish
- You'll catch a lot of what you wanted, but you'll get increasingly more bycatch

F-Measure = F_1 Measure = F Score

$$F1 = 2 * \frac{precision * recall}{precision + recall}$$

“Harmonic mean” between precision and recall

A measure that balances the trade-off between precision and recall

Range: (0,1); 1 is best

This allows us to quantify precision and recall performance with a single number

F_β Measure

$$F_\beta = (1 + \beta^2) * \frac{\text{precision} * \text{recall}}{\beta^2 * \text{precision} + \text{recall}}$$

Generalized F-Measure

You can weight the importance of precision and recall in your metric by changing β


$\beta = 2$ weights recall more

$\beta = 0.5$ weights precision more

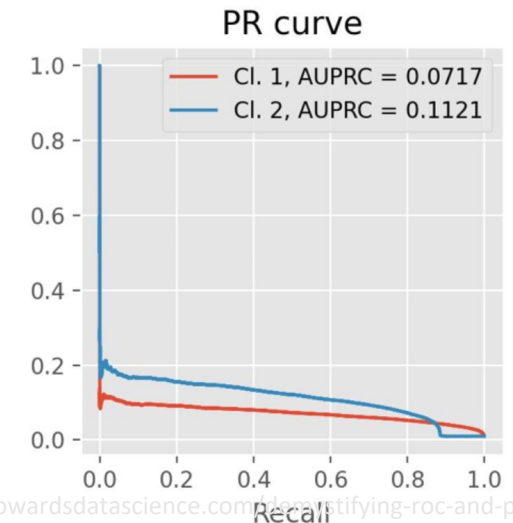
Good to know, but F1-Measure is most commonly used

Precision and Recall (PR) Curve

- Plots the trade-off between precision and recall
 - **Area under the Curve (AUC)** quantifies performance with a single number
 - Higher AUC means better, more robust system (less trade-off between precision and recall)
 - Generate by varying some threshold with some interval
 - For example, threshold of a logistic function output
 - The optimum threshold is the Euclidean distance from the perfect model
 - E.g. Euclidean distance between (1,1), and (recall, precision) of the threshold
-
- The look of PR Curves can vary a lot
 - There isn't a 1-to-1 relationship between precision and recall, and the plot can be jagged.
 - Plotting a majority class baseline is useful for interpretation
 - It will be a line at the positive class percentage



<https://analyticsindiamag.com/complete-guide-to-understanding-precision-and-recall-curves/>



<https://towardsdatascience.com/quantifying-roc-and-precision-recall-curves-d30f3fad2cbf>

Receiver Operating Characteristic (ROC) Curve

- Plots the trade off between:
 - True Positive Rate (Recall) and False Positive Rate
 - Equivalent to Sensitivity vs. 1-Specificity
 - Lots of names for the same measures:
 - see: https://en.wikipedia.org/wiki/Precision_and_recall
- Area Under the ROC Curve (AUROC) can be used to quantify performance with a single number
 - Higher AUROC means better system (less trade-off between TP rate and FP rate)
- Generate




Image from: <https://glassboxmedicine.com/2019/02/23/measuring-performance-auc-auroc/>

Use with Caution:

research* shows that **ROC Curves are misleading (overly optimistic) for problems with a high class imbalance**. For these problems, precision and recall curves are preferred

*"The Precision-Recall Plot Is More Informative than the ROC Plot When Evaluating Binary Classifiers on Imbalanced Datasets" by Saito and Rehmsmeier
<https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0118432>

Multiclass/label Evaluation

Confusion Matrix for Multiple Classes

		Predicted Class		
		Class 1	Class 2	Class 3
True Class	Class 1	4	50	1
	Class 2	10	290	65
	Class 3	0	10	2

- Show the count of samples classified as each class
- The example table tells us that there is:
 - There is confusion between classes 1 and 2
 - Most of class 3 is being classified as class 2

Multiclass/label Evaluation

- For Multiclass classification and Multilabel classification evaluation:
 - Report evaluation metrics for each class/label individually
 - Report Micro and/or Macro averaging between all classes
- Macro Averaged Precision, Recall, F1
 - All classes contribute equally
 - **Equally weights the performance of each class**
 - Tells us the average performance over all classes
- Micro Averaged Precision, Recall, F1
 - Classes with large number of samples dominate
 - **Equally weights the performance of each sample**
 - Tells us the performance over all samples

Macro Averaging

...It's just the average of the metric for each class

- Macro Averaged Accuracy

$$\frac{1}{C} \sum_i^C Accuracy_i$$

- Macro Averaged Recall

$$\frac{1}{C} \sum_i^C Recall_i$$

- Macro Averaged Precision

$$\frac{1}{C} \sum_i^C Precision_i$$

- Macro Averaged F1

$$\frac{1}{C} \sum_i^C F1_i$$

Micro Averaging

- Micro Averaged Accuracy

$$\frac{\sum_{i=1}^C TP_i}{\text{Total Number of Predictions}}$$

- Micro Averaged Precision

$$\frac{\sum_{i=1}^C TP_i}{\sum_{i=1}^C TP_i + \sum_{i=1}^C FP_i}$$

Where C = the number of classes

- Micro Averaged Recall

$$\frac{\sum_{i=1}^C TP_i}{\sum_{i=1}^C TP_i + \sum_{i=1}^C FN_i}$$

- Micro Averaged F1

$$2 * \frac{\sum_{i=1}^C precision_i * recall_i}{\sum_{i=1}^C precision_i + recall_i}$$

...This looks really complicated, but we are just considering each sample regardless of class, determining if it's a TP, FP, TN, or FN and then calculating as each metric as before

Multiclass Evaluation

		Predicted Class		
		Class 1	Class 2	Class 3
True Class	Class 1	4	50	1
	Class 2	10	290	65
	Class 3	0	10	2

- Calculate Class 1,2,3 precisions, the micro, and macro precisions
- Calculate Class 1 Recall
- Calculate Micro Recall
- Calculate Micro Accuracy

$$\text{Precision} = \frac{TP}{TP+FP}$$

$$\text{Recall} = \frac{TP}{TP+FN}$$

- Micro Averaged Precision

$$\frac{\sum_{i=1}^C TP_i}{\sum_{i=1}^C TP_i + \sum_{i=1}^C FP_i}$$

- Macro Averaged Precision

$$\frac{1}{C} \sum_i^C Precision_i$$

- Micro Averaged Accuracy

$$\frac{\sum_{i=1}^C TP_i}{\text{Total Number of Predictions}}$$

Where C = the number of classes

Multilabel Evaluation

- For multiclass problems (when each sample belongs to exactly 1 class), micro precision = micro recall = micro accuracy
- This is **not** the case for multi-label problems
 - One sample may have multiple labels, and a sample may have NO labels
- Consider the following True (Y) and Predicted (\hat{Y}) labels

$$Y = \begin{bmatrix} 1 & 1 \\ 1 & 0 \\ 0 & 0 \\ 0 & 0 \\ 0 & 1 \end{bmatrix}$$

$$\hat{Y} = \begin{bmatrix} 1 & 1 \\ 0 & 0 \\ 1 & 0 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}$$

Construct the confusion matrix
Calculate micro precision, recall, f1, accuracy

Including a “None” in the confusion matrix is useful, but it isn’t counted towards the true positives

Other Evaluation Measures

- There are a lot of evaluation measures
- Common Regression Problem evaluation measures:
 - Mean Squared Error (MSE)
 - Root-Mean-Squared-Error (RMSE)
 - Mean-Absolute-Error (MAE)
 - R^2 = Coefficient of Determination
- Common Ranking Algorithm Measures
 - Precision at K
 - Mean average precision
- Example of ranking systems are information retrieval and recommendation systems

Compare Against Baseline Systems

- Majority class baseline for classification
 - Classify everything as the majority class
- Mean baseline for regression
 - Output the mean of the data for every sample
- Can you think of another, problem specific simple classifier?
 - Compare against that
 - Sometimes a simple solution is frustratingly hard to beat

Recap:

- Different problem types:
 - Regression, binary classification, multiclass classification, multilabel classification
- How to construct and interpret a confusion matrix
 - This can help with error analysis
 - The error types can be used to derive new evaluation metrics
- How to evaluate the performance of a model using a variety of metrics
 - accuracy, F1, precision, recall, ROC curves
 - macro vs. micro measures
- We should always compare against:
 - A majority class baseline for classification, mean baseline for regression
 - and if possible against some other a simple baseline (which you implement)

Discussion

Class	Precision	Recall	F1
Micro	0.80	0.81	0.81
Macro	0.78	0.66	0.69
TrIP	0.70	0.56	0.62
TrWP	1.00	0.30	0.46
TrCP	0.75	0.79	0.77
TrAP	0.81	0.91	0.86
TrNAP	0.68	0.55	0.61
TeRP	0.86	0.88	0.87
TeCP	0.68	0.54	0.60
PIP	0.77	0.76	0.77

Relationship	TrIP	TrWP	TrCP	TrAP	TrNAP	TeRP	TeCP	PIP	Total Relations	No Relation
Total Count	203	133	526	2617	174	3053	504	2203	9413	10,308
Unique Sentence Count	140	86	370	1643	119	1858	345	1447	6008	10,308

Dealing with Imbalanced Data

- Learning on severely imbalanced datasets can be difficult
 - Classifiers can learn to just predict everything as the majority class
- This can be hard to overcome, however there are a few simple methods:
 1. Add class weights
 2. Artificially balance the dataset
 - Under-sample majority class
 - Over-sample minority class
 - Generate synthetic data for the minority class (SMOTE)

Option 1: Add Class Weights

- We can add class weights to make samples of different classes contribute more or less to loss

$$Training_Loss = J(\theta) = -\frac{1}{n} \sum_{i=1}^n w_{c,i} * loss(\hat{y}_i, y_i)$$

Scale the loss by the class weight

where $w_{c,i}$ is the class weight of the class of class i

- By **default**, the class weight balances the contribution of each class

$$w_c = 1 - \frac{\text{count of samples of class } c}{\text{total samples in the dataset}}$$

You don't have to use the default. You can select any class weights you want. Determining optimum class weights is a hyperparameter

Class Weights Example

$$w_c = 1 - \frac{\text{count of samples of class } c}{\text{total samples in the dataset}}$$

55 samples of class 1

365 samples of class 2

12 samples of class 3

432 Samples total

$$w_1 = 1 - \frac{55}{432} = 1 - 0.127 = 0.873$$

$$w_2 = 1 - \frac{365}{432} = 1 - 0.845 = 0.155$$

$$w_3 = 1 - \frac{12}{432} = 1 - 0.028 = 0.972$$

		Predicted Class		
		Class 1	Class 2	Class 3
True Class	Class 1	4	50	1
	Class 2	10	290	65
	Class 3	0	10	2

$$Training_Loss = J(\theta) = -\frac{1}{n} \sum_{i=1}^n w_{c,i} * loss(\hat{y}_i, y_i)$$

where $w_{c,i}$ is the class weight of the class of class i

So, each time we incorrectly classify a sample of class 3, it adds about 6 times more loss than missing a sample from class 2

This changes the loss surface, and simply classifying everything as class 2 may no longer be a good option (therefore removing lots of local minimum)

Option 2: Artificially Balance the Dataset

1. **Under-sample** the majority class


- Randomly select a fixed number of samples from each class such that the resulting dataset is balanced
- Problem: you are removing data which may be informative to the system

2. **Over-sample** the minority class

- Randomly repeat a fixed number of samples from each minority class such that the resulting dataset is balanced
- Problem: you are repeating data which makes those points artificially more important. What if the points are outliers?

$$\text{Training_Loss} = J(\theta) = -\frac{1}{n} \sum_{i=1}^n \text{loss}(\hat{y}_i, y_i)$$

Each of these methods make it so that each class contributes equally to loss



Option 2: Artificially Balance the Dataset

3. **Generate synthetic data** for the minority class

- This is typically a better option than over-sampling, but could be problematic if your synthetic data is not representative of your real data.
- A popular data generation method is “[Synthetic Minority Oversampling Technique](#)” (SMOTE)
- The authors recommend a combination of under-sampling the majority class and using SMOTE to over-sample the minority class. Stating that it outperforms doing only under-sampling or doing only over-sampling.

“Synthetic Minority Oversampling Technique” (SMOTE)

- SMOTE works by:
 1. For a single point, find the k-nearest neighbors and randomly select one
 2. For each feature:
 1. Find the difference between the sample and the its neighbor
 2. Randomly generate a number between 0 and 1 and multiply it by the difference
 3. Add the scaled difference to the feature value of that original point
- This effectively creates new feature values somewhere on a line between the point and its nearest neighbor.
- “This approach effectively forces the decision region of the minority class to become more general.”

Notes on Class Weights and Artificially Balancing the Dataset

- It is best, if possible to learn with the original data distribution
- **Potential Drawbacks:**
 - Class weights/oversampling can give excessive importance to outliers
 - So, for severely imbalanced datasets, you may not want to reweight/sample so that all classes are equally weighted in the loss calculation
 - Under-sampling can eliminate important datapoints
 - Particularly if the samples for a class are already sparse.
 - When you modify the dataset, you learn something that is not representative of the real world.
 - Weighting or balancing the dataset biases the system
 - If you are unsure of a sample, maybe you should classify it as the majority class rather than giving equal “weight”

Implementation Detail: Stratified Sampling

- When you have severely imbalanced data, it is important to use stratified sampling for train/validation/test splits and cross-validation
- Stratified sampling ensures a similar class distribution per split
 - E.g. if the dataset is 95% negative, and 5% positive, then each fold of cross-validation will have a 95%-5% class balance
 - This is done by sampling each class independently
- Stratified sampling is more complex for multi-label problems, and the exact solution is unclear

Discuss

(I believe there is no implementation of this in scikit-learn)

Section Summary

We learned how to deal with complex problems

1. How to evaluate performance for problems with class imbalances
 - Precision, recall, F1
2. How to evaluate performance for multiclass and multilabel problems
 - Macro vs. micro performance
3. Practical considerations for problems with severe class imbalance